

# Embedded Systems Design and Modeling



## Chapter 12 Part 2 Scheduling Anomalies

# Outline

---

- ❑ Scheduling anomalies definition
- ❑ Scheduling anomalies cases
  - Caused by mutual exclusion:
    1. Priority inversion
      - Solution: priority inheritance
      - Real case: Mars Pathfinder
    2. Deadlock
      - Solution: priority ceiling protocol
  - Seen in multiprocessor environments:
    3. Richard's anomalies (non-monotonic, brittle)
    4. Again, mutual exclusion issues

# Basics


---

- ❑ Scheduling anomalies definition: when a schedule shows unexpected and counterintuitive behaviors under special circumstances
- ❑ Often (but not always) caused by mutual exclusion locks:
  - Mutexes are needed to control accesses to shared resources
  - They can also complicate the scheduling and cause anomalies

# Anomalies Caused by Mutexes

## 1. Priority inversion:

- A high priority task is ready to execute but is blocked by a lower priority task that holds a lock it needs
- It can be bounded or unbounded
- Example: task 1 has highest priority, task 3 lowest. Task 3 acquires a lock on a shared resource. It gets preempted by task 1, which then tries to acquire the lock and blocks. Task 2 preempts task 3 at time 4, keeping the higher priority task 1 blocked for a large amount of time.



# Real Example: Mars Pathfinder

---

- ❑ The Mars Rover Pathfinder landed on Mars on July 4th, 1997.
- ❑ After a few days the Pathfinder began missing deadlines, losing data, and self-resets.
- ❑ The problem was diagnosed on the ground as priority inversion.
- ❑ Two tasks were critical for controlling communication on Pathfinder's communication bus: the scheduler task (`bc_sched`) and the distribution task (`bc_dist`).
- ❑ Each of these tasks checked every 125ms to be sure that the other had run successfully.


# Pathfinder Story (Continued)

---

- ❑ bc\_dist was blocked by a much lower priority meteorological science task (ASI/MET).
- ❑ ASI/MET was preempted by several medium priority processes such as accelerometers and radar altimeters.
- ❑ bc\_sched started and discovered that bc\_dist had not completed. Under these circumstances, bc\_sched reacted by reinitializing the lander's hardware and software and terminating all ground command activities.
- ❑ NASA and WindRiver reproduced the failure on Earth and discovered the priority inversion.

# Priority Inversion Solution


- Priority inheritance: when a task blocks the execution of another higher priority task, it executes at the highest priority of all of the tasks it blocks.
- Example: Task 1 has highest priority, task 3 lowest. Task 3 acquires a lock on a shared object, entering a critical section. It gets preempted by task 1, which then tries to acquire the lock and blocks. Task 3 inherits the priority of task 1, preventing preemption by task 2.



# Anomalies Caused by Mutexes

## 2. Deadlock:

- The lower priority task starts first and acquires lock a, then gets preempted by the higher priority task, which acquires lock b and then blocks trying to acquire lock a. The lower priority task then blocks trying to acquire lock b, and no further progress is possible.






# Deadlock Solution

---

- Priority Ceiling Protocol: Every lock or semaphore is assigned a priority ceiling equal to the priority of the highest-priority task that can potentially lock it.
- A task can acquire a lock only if the task's priority is strictly higher than the priority ceilings of all locks currently held by other tasks.
- This prevents deadlocks by blocking a task to acquire a lock held by other tasks.
- There are extensions supporting dynamic priorities and dynamic creations of locks.

# Priority Ceiling Protocol Example

- Locks a and b have priority ceilings equal to the priority of task 1. At time 3, task 1 attempts to lock b, but it cannot because task 2 currently holds lock a, which has priority ceiling equal to the priority of task 1.



# Partial Summary

---

- To successfully share resources, a system needs two properties:
  - Freedom from mutual deadlock
  - Freedom from unbounded priority inversion
    - Is bounded priority inversion acceptable?
- The combination of priority inheritance protocol and the priority ceiling protocol guarantee the above properties.

# Anomalies in Multiprocessors

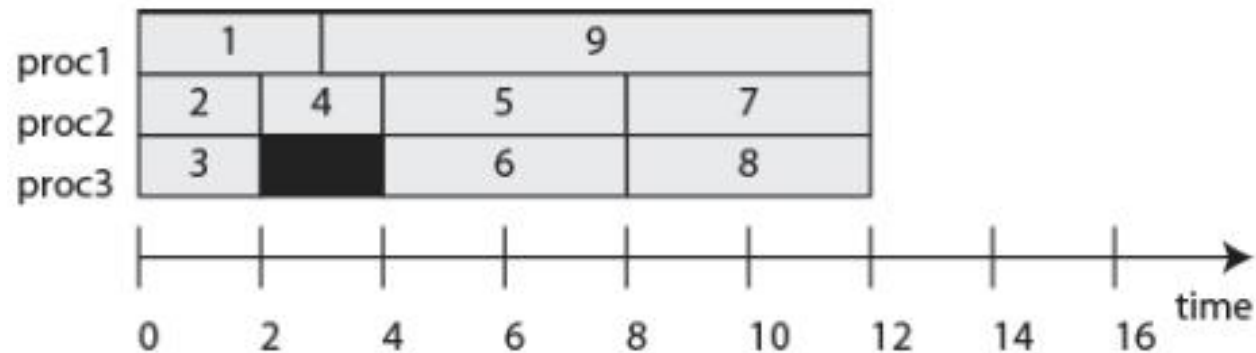
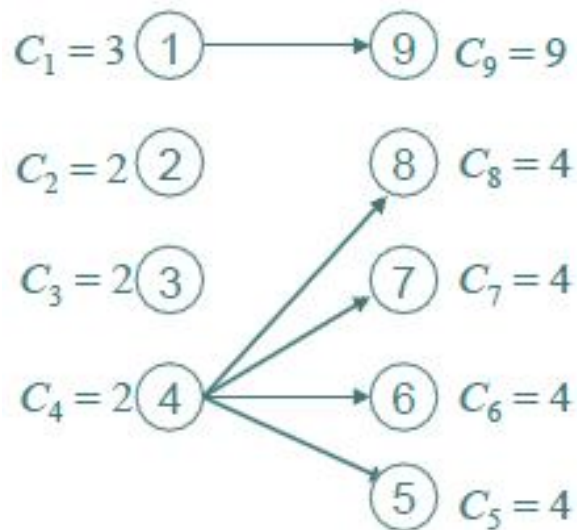
---

## 3. Known as Richard's anomalies


- Theorem: If a task set with fixed priorities, execution times, and precedence constraints is scheduled according to priorities on a fixed number of processors, then increasing the number of processors, reducing execution times, or weakening precedence constraints may not improve the schedule length and may even make it longer.

# Richard's Anomalies

- Consider 9 tasks with the following precedence graph and execution times.
- Assume lower numbered tasks have higher priority than higher numbered tasks.
- The priority-based three-processor schedule:





# Adding One More Processor



- ❑ The four-processor schedule takes longer!
- ❑ Priority-based scheduling is greedy. A smarter scheduler for this example could hold off scheduling 5, 6, or 7, leaving a processor idle for one time unit.
- ❑ But if tasks arrive only after their predecessor completes, then greedy scheduling may be the **only practical option**.

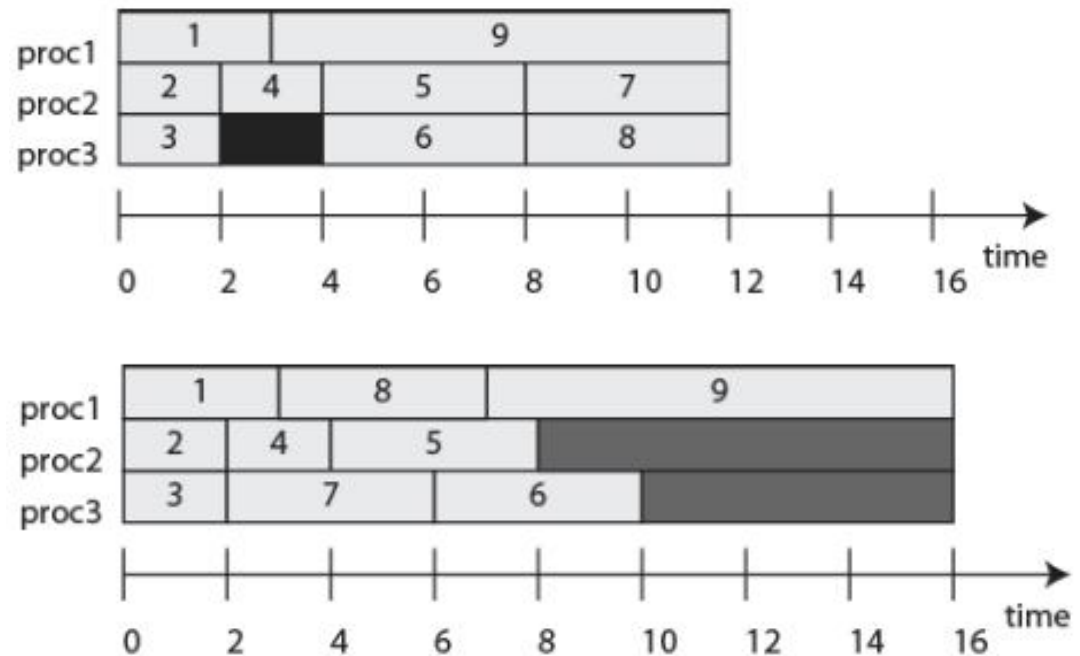

# Reducing Execution Times By 1



- Reducing the computation times by 1 also results in a longer execution time!
- Again, this is caused by the greedy approach due to dynamic scheduling.

# Weakening Precedence Constraints

- Removing the precedence constraints (4,8) and (4,7):





- Weakening precedence constraints can also result in a longer schedule!



# Anomalies in Multiprocessors

4. Anomalies caused by mutexes:
- Assume tasks 2 and 4 share the same resource in exclusive mode, and tasks are statically allocated to processors. Then if the execution time of task 1 is reduced, the schedule length increases.



# Conclusions

---

- ❑ In general, all scheduling algorithms suffer from possible anomalies.
- ❑ Timing behavior under all known task scheduling strategies is brittle:
  - Small changes can have big and unexpected consequences.
- ❑ And is non-monotonic:
  - Improvements in performance at a local level can result in degradations in performance at a global level,
- ❑ Since execution times are hard to predict, anomalies can result in system failures.
- ❑ Chapter 12 homework: 1 thru 5 for 1403/3/8