

# ReTransformer

## ReRAM-based Processing-in-Memory Architecture for Transformer Acceleration

Reza Adinepour

Amirkabir University of Technology (Tehran Polytechnic)

`adinepour@aut.ac.ir`

Computer Engineering Department

June 28, 2024

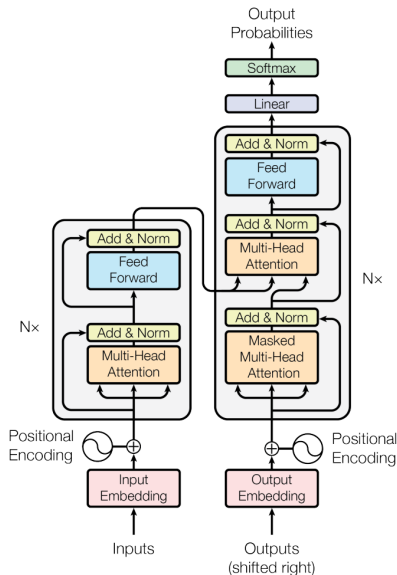


# Presentation Overview

- ➊ Introduction to Transformer Network
  - Why Use Transformers?
  - Strengths of Transformer
  - Weaknesses of Transformers
- ➋ Motivation
- ➌ ReRAM Concept
  - Vector-Matrix Multiplication (VMM)
  - Matrix-Matrix Multiplication (MatMul)
  - In-Memory Logic Operations
- ➍ Overall Architecture
  - Key Features
  - Workflow
- ➎ Experimental Setup
- ➏ Results and Analysis
- ➐ References

# What is a Transformer Network?

- 1 Introduced in "**Attention is All You Need**" (2017)
- 2 Unlike traditional RNNs and LSTMs, Transformers **do not** process data sequentially but use a mechanism called "**self-attention**" to draw global dependencies between input and output.



# Why Use Transformers?

## ① **Parallelization:**

Unlike RNNs, Transformers can process input data in parallel, leading to faster training times.

## ② **Self-Attention Mechanism:**

This allows the model to weigh the importance of different words in a sentence, capturing long-range dependencies more effectively.

## ③ **Scalability:**

Transformers can be scaled up effectively, leading to improved performance with more data and larger models.

## ④ **Versatility:**

They are used in various applications, from machine translation and text generation to image processing and more.

# Strengths of Transformer

## ① **Efficiency:**

Due to parallel processing, they train faster on large datasets.

## ② **Accuracy:**

State-of-the-art performance in many tasks, particularly in NLP.

## ③ **Flexibility:**

Applicable to a wide range of tasks beyond language, such as image and speech processing.

## ④ **Transfer Learning:**

Pre-trained models like BERT and GPT can be fine-tuned for specific tasks with relatively small amounts of data.

# Weaknesses of Transformers

## ① **Resource-Intensive:**

Require significant computational power and memory, especially for large models.

## ② **Complexity:**

More challenging to understand and implement compared to simpler models.

## ③ **Data Requirements:**

Performance often hinges on the availability of large-scale datasets for pre-training.

## ④ **Inference Speed:**

Performance bottlenecks during inference due to the scaled dot-product attention mechanism.

# Motivation

In this paper:

- ① Developed **ReTransformer**, a **ReRAM-based Processing-in-Memory (PIM) architecture** specifically designed to accelerate Transformer models.
- ② Implemented **optimized MatMul operations** to reduce data dependency and intermediate result handling
- ③ Designed a hybrid softmax mechanism combining in-memory logic and look-up tables for efficient softmax calculations.
- ④ Introduced a **sub-matrix pipeline design** for better utilization of ReRAM crossbars and improved throughput.

# Motivation (Cont.)

And Improvements Made is:

❶ **Computing Efficiency:**

23.21x improvement over GPU, 3.25x over PipeLayer.

❷ **Power Consumption:**

1086x reduction compared to GPU, 2.82x compared to PipeLayer.

❸ **Latency Reduction:**

1.32x for smaller models, 1.16x for larger models.

❹ **Softmax Efficiency:**

32% lower power consumption compared to traditional CMOS-based designs.

❺ **Throughput Enhancement:**

1.18x increase in computational throughput.



# Motivation (Cont.)

This concept use in CNNs and RNNs. but we can't and cant be directly applied to Transformer due to the following reasons:

❶ **Matrix-Matrix Multiplication:**

Transformers require frequent matrix-matrix multiplications, causing potential slowdowns and reduced efficiency due to intermediate result storage.

❷ **Different Computations:**

Unlike CNNs, Transformers use scaled dot-product attention, necessitating different computational approaches.

❸ **Finer Pipeline Granularity:**

Transformer accelerators need a more detailed pipeline design compared to the layer-level granularity used in previous designs.

# ReRAM Concept

## ReRAM Basics:

- ① Is a Non-volatile memory with:
  - High density
  - Low access energy
  - And support for multi-level cell and 3D integration
- ② ReRAM-based Vector-Matrix & Matrix-Matrix Multiplication:
  - Vector-Matrix Multiplication (VMM)
  - Matrix-Matrix Multiplication (MatMul)
- ③ In-Memory Logic Operations:
  - NOR Logic
  - XOR Logic
  - And other Logic Operations

# Vector-Matrix Multiplication (VMM)

- 1 Conductance of ReRAM cells represents elements in a matrix.

- 2 Input voltage vector ( $V_I = [v_0, v_1, v_2, v_3]$ ) fed to word lines (WLs) generates output current through bit lines (BLs).

- 3 One read cycle completes the VMM operation.

According to Kirchhoff's law the output current calculate as bellow:

$$i_j = \sum_{i=0}^3 \frac{v_i}{R(i,j)} = \sum_{i=0}^3 v_i G(i,j)$$




Figure: ReRAM-based vector-matrix multiplication

# Matrix-Matrix Multiplication (MatMul)

- ① Input matrix separated into vectors for sequential VMM operations
- ② Results of VMM operations combined to obtain MatMul results

# In-Memory Logic Operations

## ① NOR Logic:

Uses high and low conductance values to represent logic states. Operations performed using specific voltage settings.

## ② XOR Logic:

Implemented using a combination of OR and NAND operations, leveraging ReRAM's programmable conductance.

## ③ Other Logic Operations:

INV and OR implemented in one or two cycles respectively.

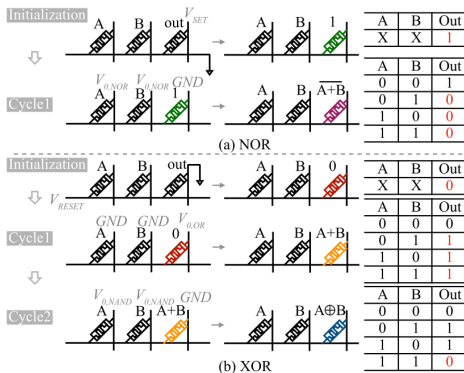


Figure: ReRAM-based in-memory logic: (a) NOR, (b) XOR.

# Overall Architecture

A ReRAM-based PIM module is divided into three types of functional components:

## ① Processing Subarrays:

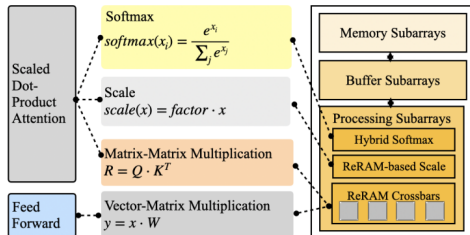
Execute computations such as MatMul and feed-forward operations.

## ② Buffer Subarrays:

Serve as caches to store intermediate data and results.

## ③ Memory Subarrays:

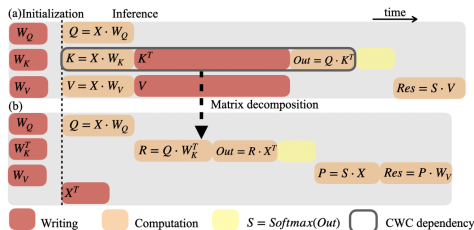
Store original input data and final output results.



**Figure:** Overview of the proposed ReRAM-based PIM design for Transformer

# Key Features

- 1 **ReRAM Crossbars:**  
Enable efficient in-memory computations.
- 2 **Optimized MatMul:**  
Reduces data dependency and intermediate writes.
- 3 **Hybrid Softmax Mechanism:**  
Combines in-memory logic with look-up tables for efficiency.
- 4 **Sub-Matrix Pipeline:**  
Slices input matrices for parallel processing.



**Figure:** Remove data dependency in scaled dot-product attention layer: (a) a CWC dependency caused by the in- intermediate result  $K$ . (b) The optimized MatMul eliminates the CWC dependency by decomposing the computation into two cascaded multiplications.

## ① Data Loading:

Input data is stored in memory subarrays.

## ② Computation:

Processing subarrays perform operations using data from buffer subarrays.

## ③ Intermediate Handling:

Buffer subarrays temporarily store intermediate results, reducing the need for frequent memory writes.

## ④ Output Storage:

Final results are stored back in memory subarrays.



# Experimental Setup

## ① Configurations:

- GPU: NVIDIA TITAN RTX, 24GB memory, 672 GB/s bandwidth.
- ReTransformer: ReRAM crossbar arrays, 2-bit cell precision, 128x128 subarray size.

## ② Evaluation Metrics:

- Computing Efficiency: Operations per second per watt.
- Power Consumption: Total power used during computations.

# Results and Analysis

## MatMul Optimization:

Reduces computation latency by 1.32x for Model A and 1.16x for Model B.

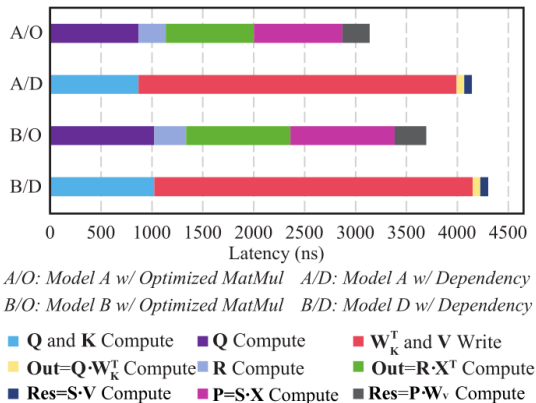


Figure: MatMul computation latency comparisons

# Results and Analysis (Cont.)

## Hybrid Softmax Efficiency:

Lowers power consumption by 32% compared to CMOS-based design.

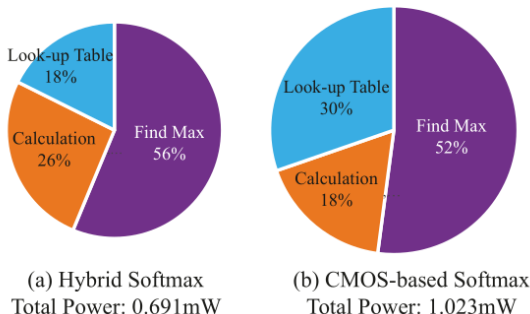


Figure: Softmax design comparison

# Results and Analysis (Cont.)

## Pipeline Performance:

Finer granularity pipeline improves throughput by 1.18x for both models.

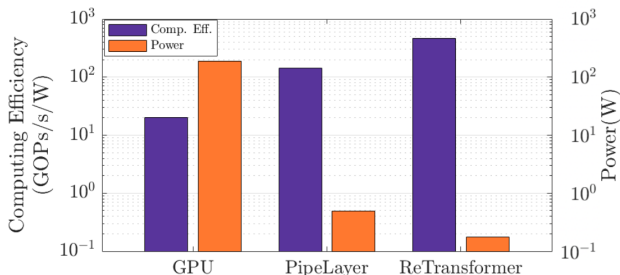
Model	Layer	Finer	Improvement
Model A	69.24 GOPs/s	81.85 GOPs/s	1.18x
Model B	67.89 GOPs/s	80.07 GOPs/s	1.18x

Table: Performance comparison of two pipeline designs

# Results and Analysis (Cont.)

## Overall Comparison:

- ① ReTransformer achieves 23.21x improvement in computing efficiency and 1086x reduction in power consumption compared to GPU.
- ② Compared to PipeLayer, ReTransformer improves computing efficiency by 3.25x and reduces power by 2.82x.



**Figure:** Performance comparison with GPU and PipeLayer

# References

◀ Back to start



X. Yang, B. Yan, H. Li, and Y. Chen, “Retransformer: Reram-based processing-in-memory architecture for transformer acceleration,” in *Proceedings of the 39th International Conference on Computer-Aided Design*, 2020, pp. 1–9.



A. Vaswani, N. Shazeer, N. Parmar, J. Uszkoreit, L. Jones, A. N. Gomez, Ł. Kaiser, and I. Polosukhin, “Attention is all you need,” *Advances in neural information processing systems*, vol. 30, 2017.

# The End

## Questions? Comments?

You can find this slides here:

[github.com/M-Sc-AUT/M.Sc-Computer-Architecture/MemoryTechnologies](https://github.com/M-Sc-AUT/M.Sc-Computer-Architecture/MemoryTechnologies)

I delivered three presentations this semester on the following topics:

- ❶ Multiprocessors shared-memory Architecture
- ❷ Exploring CACTI and NVSIM Simulators
- ❸ Exploring DRAMSim and SimpleScaler Simulators