

ReTransformer

ReRAM-based Processing-in-Memory Architecture for Transformer Acceleration

Reza Adinepour

Amirkabir University of Technology (Tehran Polytechnic)

`adinepour@aut.ac.ir`

Computer Engineering Department

June 28, 2024

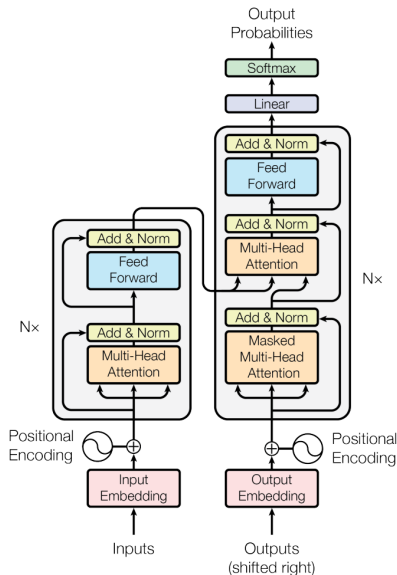


Presentation Overview

- ➊ Introduction to Transformer Network
 - Why Use Transformers?
 - Strengths of Transformer
 - Weaknesses of Transformers
- ➋ Motivation
- ➌ ReRAM Concept
 - Vector-Matrix Multiplication (VMM)
 - Matrix-Matrix Multiplication (MatMul)
 - In-Memory Logic Operations
- ➍ Overall Architecture
 - Key Features
 - Workflow
- ➎ Experimental Setup
- ➏ Results and Analysis
- ➐ References

What is a Transformer Network?

- 1 Introduced in "**Attention is All You Need**" (2017)
- 2 Unlike traditional RNNs and LSTMs, Transformers **do not** process data sequentially but use a mechanism called "**self-attention**" to draw global dependencies between input and output.



Why Use Transformers?

① **Parallelization:**

Unlike RNNs, Transformers can process input data in parallel, leading to faster training times.

② **Self-Attention Mechanism:**

This allows the model to weigh the importance of different words in a sentence, capturing long-range dependencies more effectively.

③ **Scalability:**

Transformers can be scaled up effectively, leading to improved performance with more data and larger models.

④ **Versatility:**

They are used in various applications, from machine translation and text generation to image processing and more.

Strengths of Transformer

① **Efficiency:**

Due to parallel processing, they train faster on large datasets.

② **Accuracy:**

State-of-the-art performance in many tasks, particularly in NLP.

③ **Flexibility:**

Applicable to a wide range of tasks beyond language, such as image and speech processing.

④ **Transfer Learning:**

Pre-trained models like BERT and GPT can be fine-tuned for specific tasks with relatively small amounts of data.

Weaknesses of Transformers

① **Resource-Intensive:**

Require significant computational power and memory, especially for large models.

② **Complexity:**

More challenging to understand and implement compared to simpler models.

③ **Data Requirements:**

Performance often hinges on the availability of large-scale datasets for pre-training.

④ **Inference Speed:**

Performance bottlenecks during inference due to the scaled dot-product attention mechanism.

Motivation

In this paper:

- ① Developed **ReTransformer**, a **ReRAM-based Processing-in-Memory (PIM) architecture** specifically designed to accelerate Transformer models.
- ② Implemented **optimized MatMul operations** to reduce data dependency and intermediate result handling
- ③ Designed a hybrid softmax mechanism combining in-memory logic and look-up tables for efficient softmax calculations.
- ④ Introduced a **sub-matrix pipeline design** for better utilization of ReRAM crossbars and improved throughput.

Motivation (Cont.)

And Improvements Made is:

❶ **Computing Efficiency:**

23.21x improvement over GPU, 3.25x over PipeLayer.

❷ **Power Consumption:**

1086x reduction compared to GPU, 2.82x compared to PipeLayer.

❸ **Latency Reduction:**

1.32x for smaller models, 1.16x for larger models.

❹ **Softmax Efficiency:**

32% lower power consumption compared to traditional CMOS-based designs.

❺ **Throughput Enhancement:**

1.18x increase in computational throughput.

Motivation (Cont.)

This concept use in CNNs and RNNs. but we can't and cant be directly applied to Transformer due to the following reasons:

❶ **Matrix-Matrix Multiplication:**

Transformers require frequent matrix-matrix multiplications, causing potential slowdowns and reduced efficiency due to intermediate result storage.

❷ **Different Computations:**

Unlike CNNs, Transformers use scaled dot-product attention, necessitating different computational approaches.

❸ **Finer Pipeline Granularity:**

Transformer accelerators need a more detailed pipeline design compared to the layer-level granularity used in previous designs.

ReRAM Concept

ReRAM Basics:

- ① Is a Non-volatile memory with:
 - High density
 - Low access energy
 - And support for multi-level cell and 3D integration
- ② ReRAM-based Vector-Matrix & Matrix-Matrix Multiplication:
 - Vector-Matrix Multiplication (VMM)
 - Matrix-Matrix Multiplication (MatMul)
- ③ In-Memory Logic Operations:
 - NOR Logic
 - XOR Logic
 - And other Logic Operations

Vector-Matrix Multiplication (VMM)

- 1 Conductance of ReRAM cells represents elements in a matrix.

- 2 Input voltage vector ($V_I = [v_0, v_1, v_2, v_3]$) fed to word lines (WLs) generates output current through bit lines (BLs).

- 3 One read cycle completes the VMM operation.

According to Kirchhoff's law the output current calculate as bellow:

$$i_j = \sum_{i=0}^3 \frac{v_i}{R(i,j)} = \sum_{i=0}^3 v_i G(i,j)$$

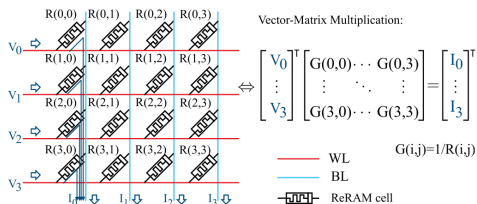


Figure: ReRAM-based vector-matrix multiplication

Matrix-Matrix Multiplication (MatMul)

- ① Input matrix separated into vectors for sequential VMM operations
- ② Results of VMM operations combined to obtain MatMul results

In-Memory Logic Operations

① NOR Logic:

Uses high and low conductance values to represent logic states. Operations performed using specific voltage settings.

② XOR Logic:

Implemented using a combination of OR and NAND operations, leveraging ReRAM's programmable conductance.

③ Other Logic Operations:

INV and OR implemented in one or two cycles respectively.

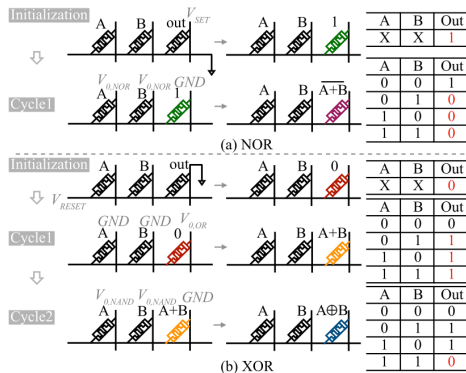


Figure: ReRAM-based in-memory logic: (a) NOR, (b) XOR.

Overall Architecture

A ReRAM-based PIM module is divided into three types of functional components:

① Processing Subarrays:

Execute computations such as MatMul and feed-forward operations.

② Buffer Subarrays:

Serve as caches to store intermediate data and results.

③ Memory Subarrays:

Store original input data and final output results.

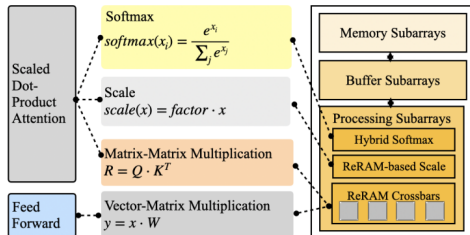


Figure: Overview of the proposed ReRAM-based PIM design for Transformer

Key Features

- 1 **ReRAM Crossbars:**
Enable efficient in-memory computations.
- 2 **Optimized MatMul:**
Reduces data dependency and intermediate writes.
- 3 **Hybrid Softmax Mechanism:**
Combines in-memory logic with look-up tables for efficiency.
- 4 **Sub-Matrix Pipeline:**
Slices input matrices for parallel processing.

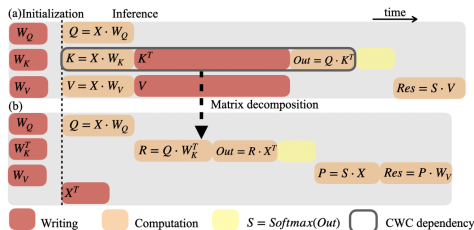


Figure: Remove data dependency in scaled dot-product attention layer: (a) a CWC dependency caused by the in- intermediate result K . (b) The optimized MatMul eliminates the CWC dependency by decomposing the computation into two cascaded multiplications.

① Data Loading:

Input data is stored in memory subarrays.

② Computation:

Processing subarrays perform operations using data from buffer subarrays.

③ Intermediate Handling:

Buffer subarrays temporarily store intermediate results, reducing the need for frequent memory writes.

④ Output Storage:

Final results are stored back in memory subarrays.

Experimental Setup

① Configurations:

- GPU: NVIDIA TITAN RTX, 24GB memory, 672 GB/s bandwidth.
- ReTransformer: ReRAM crossbar arrays, 2-bit cell precision, 128x128 subarray size.

② Evaluation Metrics:

- Computing Efficiency: Operations per second per watt.
- Power Consumption: Total power used during computations.

Results and Analysis

MatMul Optimization:

Reduces computation latency by 1.32x for Model A and 1.16x for Model B.

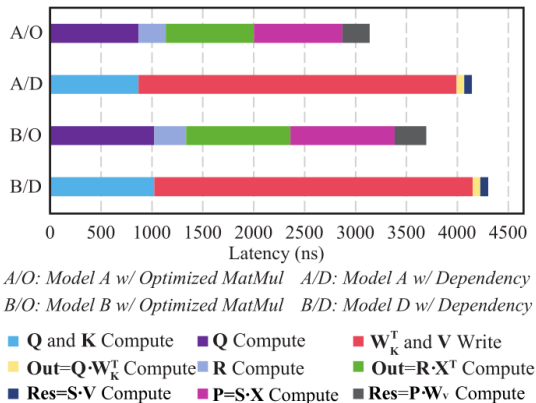


Figure: MatMul computation latency comparisons

Results and Analysis (Cont.)

Hybrid Softmax Efficiency:

Lowers power consumption by 32% compared to CMOS-based design.

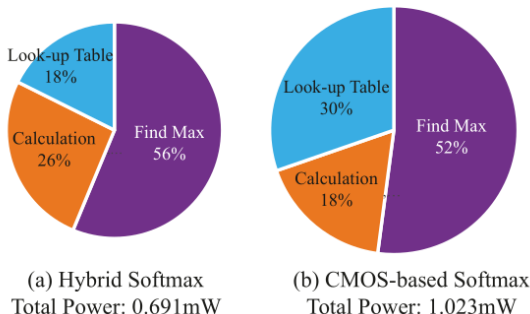


Figure: Softmax design comparison

Results and Analysis (Cont.)

Pipeline Performance:

Finer granularity pipeline improves throughput by 1.18x for both models.

| Model | Layer | Finer | Improvement |
|---------|--------------|--------------|-------------|
| Model A | 69.24 GOPs/s | 81.85 GOPs/s | 1.18x |
| Model B | 67.89 GOPs/s | 80.07 GOPs/s | 1.18x |

Table: Performance comparison of two pipeline designs

Results and Analysis (Cont.)

Overall Comparison:

- ① ReTransformer achieves 23.21x improvement in computing efficiency and 1086x reduction in power consumption compared to GPU.
- ② Compared to PipeLayer, ReTransformer improves computing efficiency by 3.25x and reduces power by 2.82x.

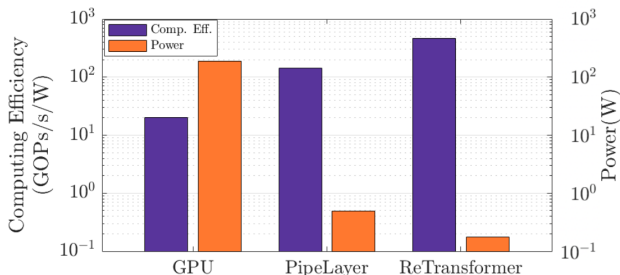


Figure: Performance comparison with GPU and PipeLayer

References

◀ Back to start



X. Yang, B. Yan, H. Li, and Y. Chen, “Retransformer: Reram-based processing-in-memory architecture for transformer acceleration,” in *Proceedings of the 39th International Conference on Computer-Aided Design*, 2020, pp. 1–9.



A. Vaswani, N. Shazeer, N. Parmar, J. Uszkoreit, L. Jones, A. N. Gomez, Ł. Kaiser, and I. Polosukhin, “Attention is all you need,” *Advances in neural information processing systems*, vol. 30, 2017.

The End

Questions? Comments?

You can find this slides here:

github.com/M-Sc-AUT/M.Sc-Computer-Architecture/MemoryTechnologies

In this term, I do 3 presentation with this topics:

- ① Multiprocessors shared-memory Architecture
- ② Exploring CACTI and NVSIM Simulators
- ③ Exploring DRAMSim and SimpleScaler Simulators