

ReTransformer

ReRAM-based Processing-in-Memory Architecture for Transformer Acceleration

Reza Adinepour

Amirkabir University of Technology (Tehran Polytechnic)

`adinepour@aut.ac.ir`

Computer Engineering Department

June 27, 2024

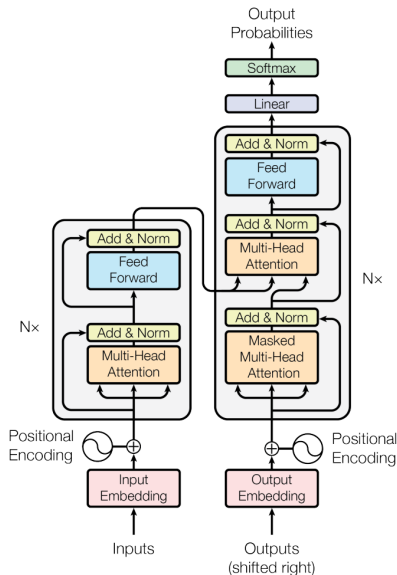


Presentation Overview

- ➊ Introduction to Transformer Network
 - Why Use Transformers?
 - Strengths of Transformer
 - Weaknesses of Transformers
- ➋ Motivation
 - Cache-only Memory Access (COMA)
- ➌ Code Snippets
- ➍ Differences between UMA and NUMA

What is a Transformer Network?

- 1 Introduced in "**Attention is All You Need**" (2017)
- 2 Unlike traditional RNNs and LSTMs, Transformers **do not** process data sequentially but use a mechanism called "**self-attention**" to draw global dependencies between input and output.



Why Use Transformers?

① **Parallelization:**

Unlike RNNs, Transformers can process input data in parallel, leading to faster training times.

② **Self-Attention Mechanism:**

This allows the model to weigh the importance of different words in a sentence, capturing long-range dependencies more effectively.

③ **Scalability:**

Transformers can be scaled up effectively, leading to improved performance with more data and larger models.

④ **Versatility:**

They are used in various applications, from machine translation and text generation to image processing and more.

Strengths of Transformer

① **Efficiency:**

Due to parallel processing, they train faster on large datasets.

② **Accuracy:**

State-of-the-art performance in many tasks, particularly in NLP.

③ **Flexibility:**

Applicable to a wide range of tasks beyond language, such as image and speech processing.

④ **Transfer Learning:**

Pre-trained models like BERT and GPT can be fine-tuned for specific tasks with relatively small amounts of data.

Weaknesses of Transformers

① **Resource-Intensive:**

Require significant computational power and memory, especially for large models.

② **Complexity:**

More challenging to understand and implement compared to simpler models.

③ **Data Requirements:**

Performance often hinges on the availability of large-scale datasets for pre-training.

④ **Inference Speed:**

Performance bottlenecks during inference due to the scaled dot-product attention mechanism.

Motivation

In this paper:

- ① Developed **ReTransformer**, a **ReRAM-based Processing-in-Memory (PIM) architecture** specifically designed to accelerate Transformer models.
- ② Implemented **optimized MatMul operations** to reduce data dependency and intermediate result handling
- ③ Designed a hybrid softmax mechanism combining in-memory logic and look-up tables for efficient softmax calculations.
- ④ Introduced a **sub-matrix pipeline design** for better utilization of ReRAM crossbars and improved throughput.

Motivation

And Improvements Made is:

❶ **Computing Efficiency:**

23.21x improvement over GPU, 3.25x over PipeLayer.

❷ **Power Consumption:**

1086x reduction compared to GPU, 2.82x compared to PipeLayer.

❸ **Latency Reduction:**

1.32x for smaller models, 1.16x for larger models.

❹ **Softmax Efficiency:**

32% lower power consumption compared to traditional CMOS-based designs.

❺ **Throughput Enhancement:**

1.18x increase in computational throughput.

Motivation

This concept use in CNNs and RNNs. but we can't and cant be directly applied to Transformer due to the following reasons:

❶ **Matrix-Matrix Multiplication:**

Transformers require frequent matrix-matrix multiplications, causing potential slowdowns and reduced efficiency due to intermediate result storage.

❷ **Different Computations:**

Unlike CNNs, Transformers use scaled dot-product attention, necessitating different computational approaches.

❸ **Finer Pipeline Granularity:**

Transformer accelerators need a more detailed pipeline design compared to the layer-level granularity used in previous designs.

Non-uniform Memory Access (NUMA) (Cont.)

Advantages:

- ① **Higher Implementation Complexity:** Additional hardware and software complexity.
- ② **Higher Latency for Remote Access:** Accessing remote memory incurs higher latency.

Example System:

① **Multi-Socket Server:**

- ① Each socket has its processors and memory banks.
- ② Sockets connected via a high-speed interconnect.
- ③ Commonly used in data centers, offers better performance for specific workloads.

Code Snippets

Run the [code!](#)

— This code is generated by ChatGP

Differences between UMA and NUMA

① Memory access time:

① NUMA:

Memory access time varies depending on the location of the data in memory. Accessing data in the local memory of a processor is faster than accessing data in the memory of a remote processor.

② UMA:

Memory access time is uniform across all processors since they share the same memory pool.

② Scalability:

① NUMA architecture is highly scalable and can support a large number of processors.

② UMA architecture is not as scalable as NUMA and may face performance issues when used with a large number of processors.

The End

Questions? Comments?

You can find this slides here:

[github.com/M-Sc-AUT/M.Sc-Computer-Architecture/Memory
Technologies](https://github.com/M-Sc-AUT/M.Sc-Computer-Architecture/MemoryTechnologies)