



# **Operating Systems**

## **System Calls**

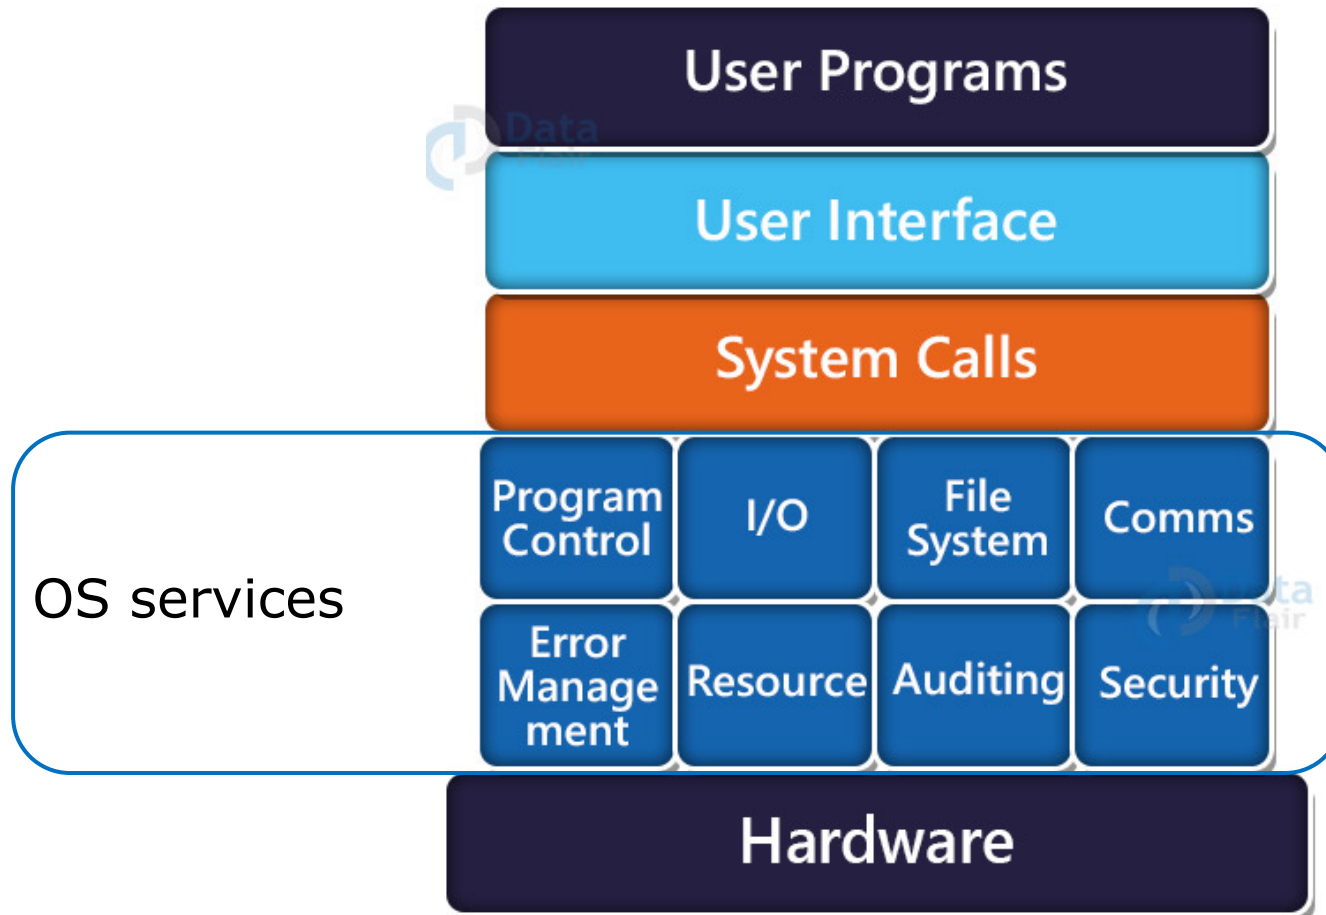
Seyyed Ahmad Javadi

[sajavadi@aut.ac.ir](mailto:sajavadi@aut.ac.ir)

Fall 2023

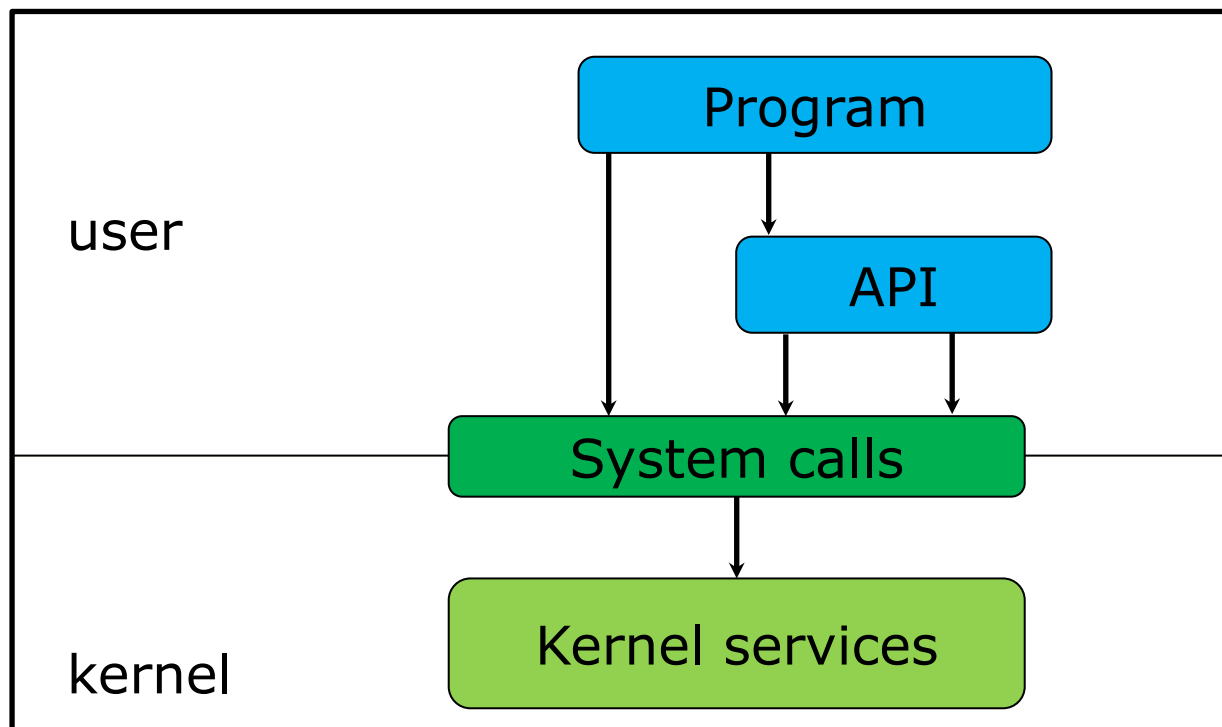
# System Calls

- **Programming interface** to the **services** provided by the OS.



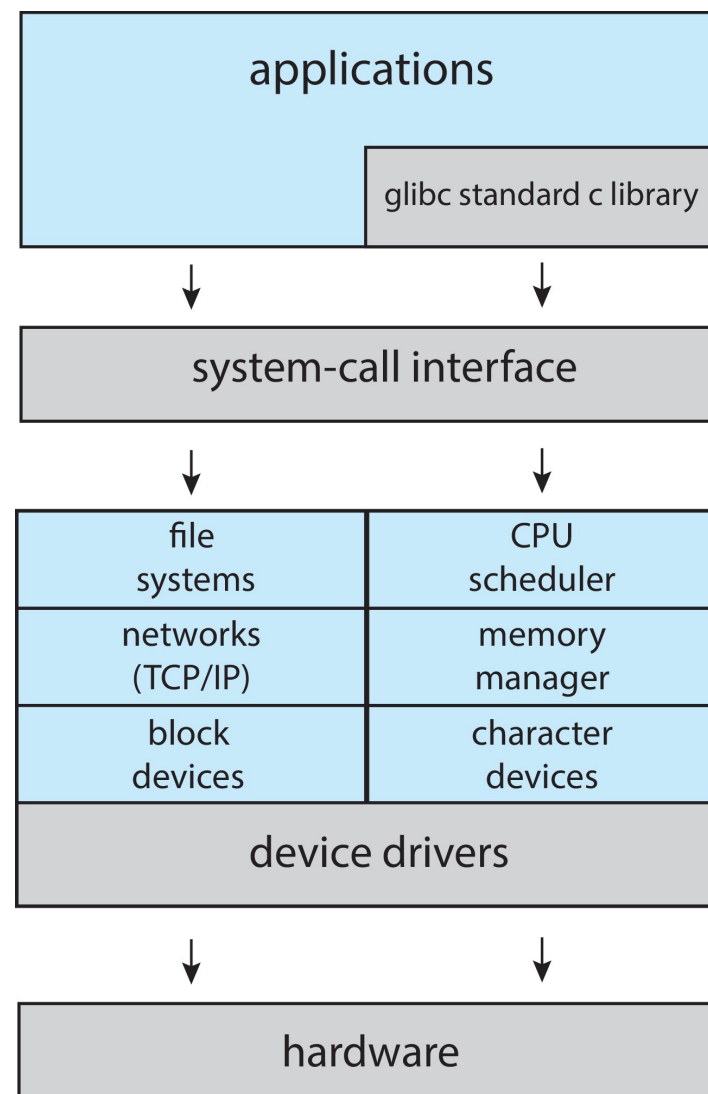
# System Calls (cont.)

- Typically written in a high-level language (C or C++ or Assembly).
- Mostly accessed by programs via a high-level **Application Programming Interface (API)** rather than direct system call use.



# Three most common APIs

- **Win32** API for Windows (Win API)
- **POSIX** API for POSIX-based systems
  - Including virtually all versions of UNIX, Linux (***unistd.h***), and Mac OS X
- **Java** API for the Java virtual machine.



# Example of Standard API

## *EXAMPLE OF STANDARD API*

As an example of a standard API, consider the `read()` function that is available in UNIX and Linux systems. The API for this function is obtained from the man page by invoking the command

```
man read
```

on the command line. A description of this API appears below:

```
#include <unistd.h>
```

```
ssize_t
```

```
read(int fd, void *buf, size_t count)
```

return  
value

function  
name

parameters

# Example of Standard API (cont.)

---

A program that uses the `read()` function must include the `unistd.h` header file, as this file defines the `ssize_t` and `size_t` data types (among other things). The parameters passed to `read()` are as follows:

- `int fd`—the file descriptor to be read
- `void *buf`—a buffer into which the data will be read
- `size_t count`—the maximum number of bytes to be read into the buffer

On a successful read, the number of bytes read is returned. A return value of 0 indicates end of file. If an error occurs, `read()` returns `-1`.



# System Call Implementation

---

- Typically, a number is associated with each system call
  - **System-call interface** maintains a table indexed according to these numbers.
- The system call interface invokes the intended system call in OS kernel and **returns status of the system call** and any **return values**



# System Call Implementation (cont.)

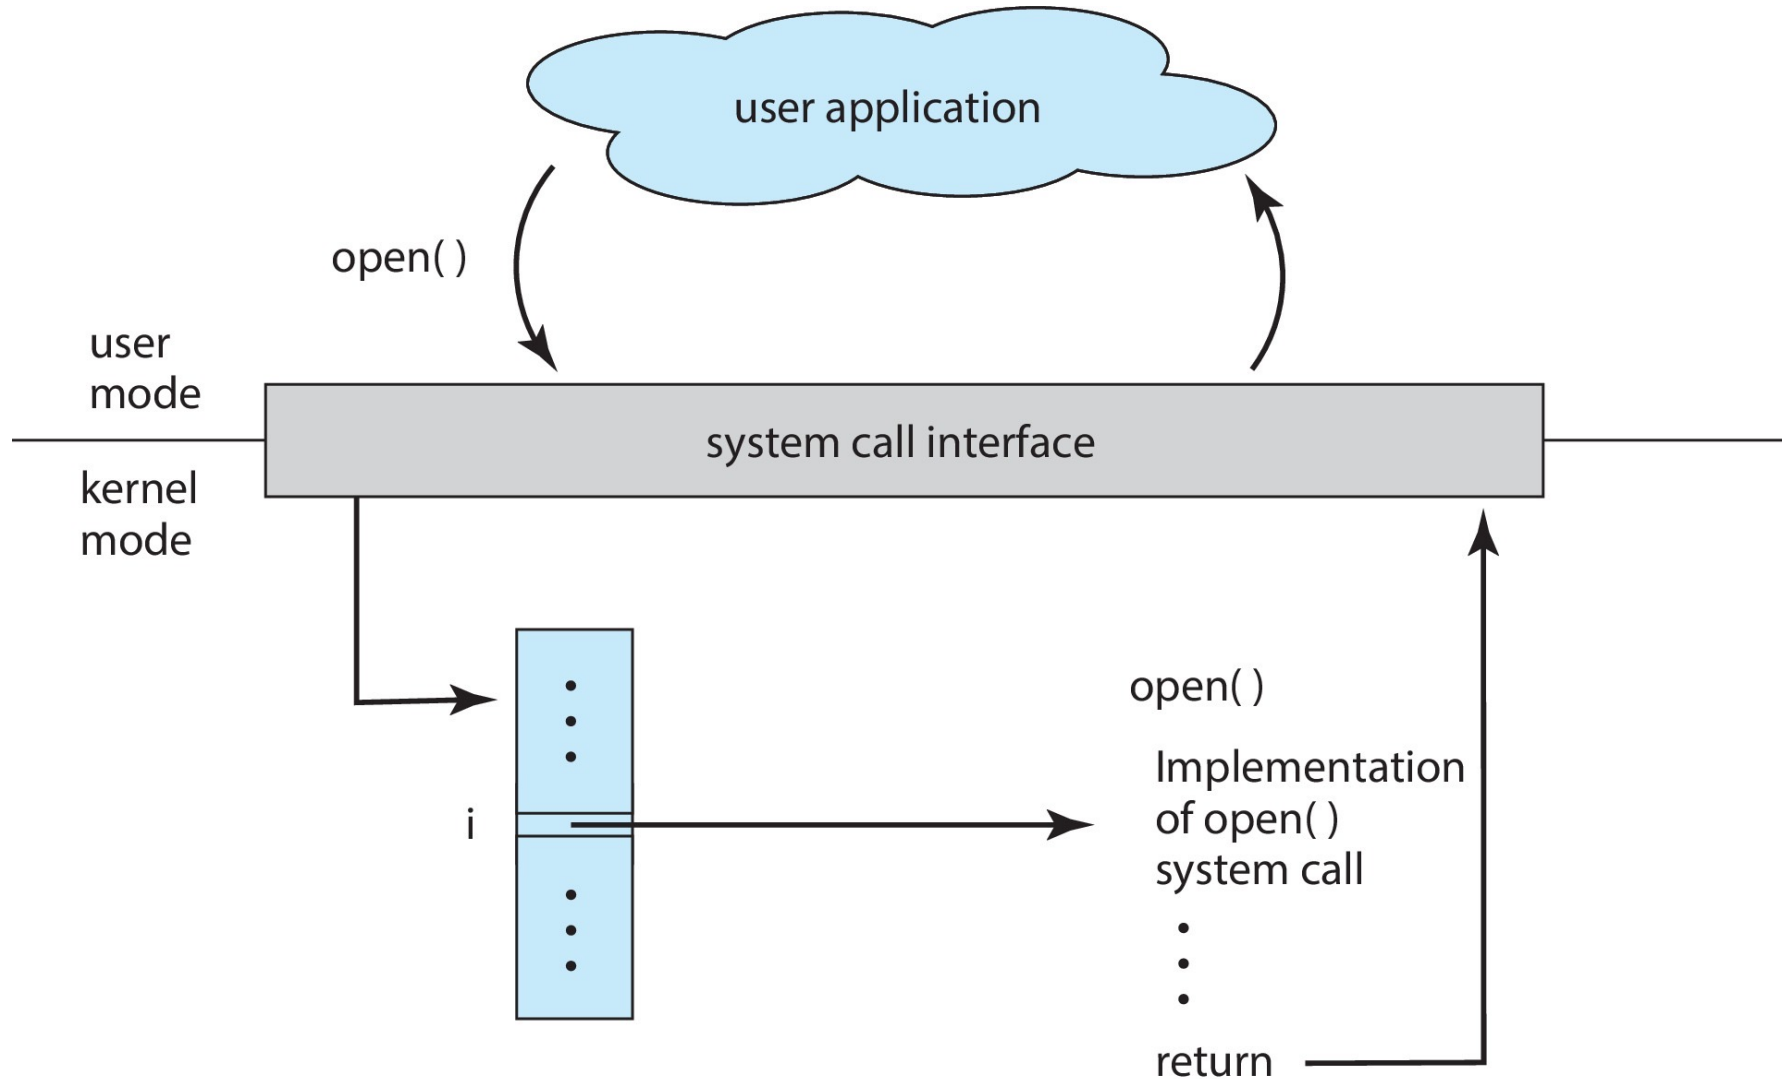
---

- The caller need know nothing about how the system call is implemented
  - Just needs to obey API and understand what OS will do as a result call.
  - Most details of OS interface hidden from programmer by API
    - ▶ Managed by run-time support library





# API – System Call – OS Relationship



# System calls in assembly programs (demo)

---

- Put the system call number in the ***EAX register***.
- Store the arguments to the system call in the registers EBX, ECX,...
- Call the relevant interrupt (***80h***).
- The result is ***usually*** returned in the ***EAX*** register.

[https://www.tutorialspoint.com/assembly\\_programming/assembly\\_system\\_calls.htm](https://www.tutorialspoint.com/assembly_programming/assembly_system_calls.htm)

[http://faculty.nps.edu/cseagle/assembly/sys\\_call.html](http://faculty.nps.edu/cseagle/assembly/sys_call.html)



# System Call Parameter Passing


---

- **Parameter Passing**
  - Register
  - Register pointer to mem. block
  - Stack (Push, Pop)
- Often, more information is required than simply identity of desired system call.
- Exact type and amount of information vary according to OS and call.



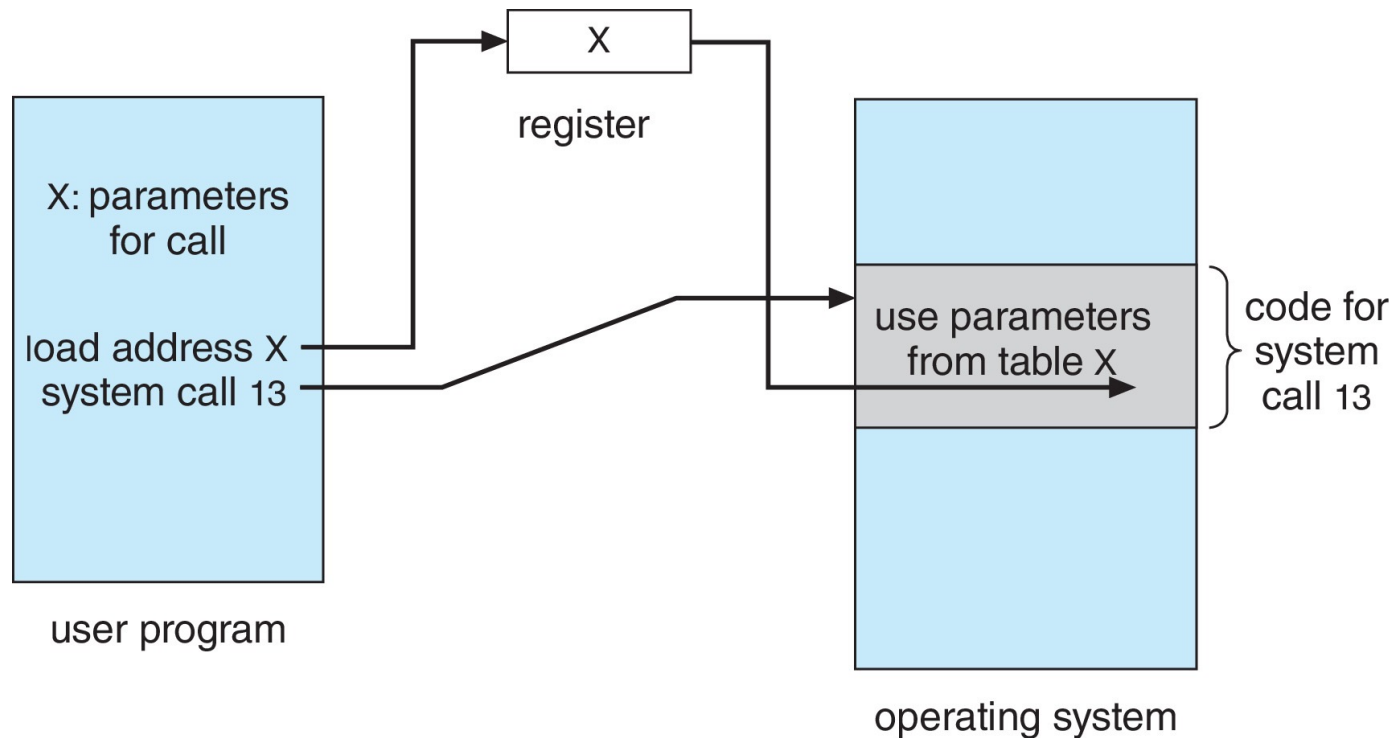
# System Call Parameter Passing--Methods

- **Simplest:** pass the parameters in registers.
  - In some cases, may be more parameters than registers.




# System Call Parameter Passing—Methods (cont.)

- Parameters stored in a **block, or table, in memory**, and address of block passed as a parameter in a register.
- This approach taken by Linux and Solaris.



# System Call Parameter Passing—Methods (cont.)

- Parameters placed, or **pushed**, onto the **stack** by the program and **popped** off the stack by the operating system.



# System Call Parameter Passing—Methods (cont.)

---

Methods\features	Is there a limitation on the number of parameters?	Is there a limitation on the length of the parameters?
<b>Register</b>		
<b>Register pointer to mem. block</b>		
<b>Stack (Push, Pop)</b>		

# System Call Parameter Passing—Methods (cont.)

---

Methods\features	Is there a limitation on the number of parameters?	Is there a limitation on the length of the parameters?
<b>Register</b>	YES	YES
<b>Register pointer to mem. block</b>	NO	NO
<b>Stack (Push, Pop)</b>	NO	NO





# Types of System Calls

---

- **Process control**
  - Create process, terminate process
  - ...
- **File management**
  - create file, delete file
  - ...
- **Device management**
  - request device, release device
  - ...
- **Please study the reference book for more details**

# Types of System Calls (Cont.)

	Windows	Unix
<b>Process Control</b>	CreateProcess() ExitProcess() WaitForSingleObject()	fork() exit() wait()
<b>File Manipulation</b>	CreateFile() ReadFile() WriteFile() CloseHandle()	open() read() write() close()
<b>Device Manipulation</b>	SetConsoleMode() ReadConsole() WriteConsole()	ioctl() read() write()
<b>Information Maintenance</b>	GetCurrentProcessID() SetTimer() Sleep()	getpid() alarm() sleep()
<b>Communication</b>	CreatePipe() CreateFileMapping() MapViewOfFile()	pipe() shm_open() mmap()
<b>Protection</b>	SetFileSecurity() InitializeSecurityDescriptor() SetSecurityDescriptorGroup()	chmod() umask() chown()



# Why Applications are Operating System Specific

---

- Apps compiled on one system usually not executable on other OSs.
- Each OS provides its own unique system calls
  - Own file formats, etc.
- Apps can be multi-operating system
  - Written in interpreted language like Python, Ruby, and interpreter available on multiple OSs.
  - App written in language that includes a VM containing the running app (like Java).
  - Use standard language (like C), compile separately on each operating system to run on each.

